AI MUSTAQBIL 2.0

Evaluation Criteria & Judging Rubric

Multi-Agentic Systems Track

> All teams in AI Mustaqbil 2.0 are evaluated on the same six criteria below. Each criterion is designed specifically for multi-agentic system challenges and applies universally regardless of the problem domain chosen.

Overview: Scoring at a Glance

Criterion	Points	What Judges Look For
Problem Relevance & Impact	**20**	Does the problem justify a multi-agent approach? Is the scope clear and impact meaningful?
Agent Design & Orchestration	**20**	Are agent roles distinct? Is the orchestration logic coherent and communication reliable?
Technical Execution	**20**	Does the pipeline run end-to-end? Are agent loops stable and failure modes handled?
Creativity & Originality	**15**	Does the architecture push beyond standard patterns? Is the use case or approach novel?
Presentation & Communication	**15**	Can the team explain agent roles clearly? Does the demo show real agent interactions?
Responsible AI & Safety	**10**	Are guardrails in place? Are irreversible actions confirmed? Is output explainable?
Total Score	**100**	Sum of all six criteria

Detailed Rubric

The following rubric gives judges precise descriptors for each score band. Use these to calibrate scores consistently across teams.

1. Problem Relevance & Impact

Score	Excellent	Satisfactory	Needs Improvement

0 – 20	High-impact, well-scoped problem that clearly requires a multi-agent approach. Impact is articulated compellingly.	Problem is real but loosely scoped, or the multi-agent justification is thin.	Problem is trivial, unclear, or a single-agent solution would suffice just as well.

2. Agent Design & Orchestration

Score	Excellent	Satisfactory	Needs Improvement
0 – 20	Agents have non-overlapping roles; orchestration strategy is clearly chosen (pipeline / hierarchical / event-driven); inter-agent communication is structured.	Agents exist but roles overlap, or orchestration is ad-hoc without a clear strategy.	Agents are superficial wrappers around a single LLM call; no real orchestration logic.

3. Technical Execution

Score	Excellent	Satisfactory	Needs Improvement
0 – 20	Full pipeline runs without manual intervention; tool use and context-passing work correctly; agent loops terminate reliably.	System mostly works but has minor bugs, requires manual resets, or some steps are mocked.	Pipeline is broken, incomplete, or the demo is mostly a slide deck with no live system.

4. Creativity & Originality

Score	Excellent	Satisfactory	Needs Improvement
0 – 15	Genuinely novel agent architecture or application domain. Creative use of memory, planning, or multi-step reasoning.	Familiar idea executed well; limited novelty in agent design.	Standard demo (e.g. basic Q&A bot) with no meaningful innovation.

5. Presentation & Communication

Score	Excellent	Satisfactory	Needs Improvement
0 – 15	Architecture diagram shown; each agent role explained simply; live demo shows real agent-to-agent interactions.	Presentation is understandable but lacks architectural clarity or relies on static slides.	Presentation is confusing; judges cannot follow what the agents are doing or why.

6. Responsible AI & Safety

Score	Excellent	Satisfactory	Needs Improvement
0 – 10	Agents confirm before irreversible actions; prompt injection considered; output is transparent and explainable to end users.	Some safety thinking present but guardrails are incomplete or inconsistently applied.	No safety considerations; agents can take unconstrained actions with no oversight.

Notes for Judges

- Score each criterion independently — do not let overall impression inflate or deflate individual scores.
- A working demo is required for full marks in Technical Execution. Mocked outputs or slide-only presentations should score in the lower band.
- Responsible AI is non-optional in multi-agent systems. Agents that take unconstrained real-world actions (send emails, call APIs, write files) with no guardrails should lose marks.
- If two teams are tied, use Creativity & Originality as the primary tiebreaker, followed by Agent Design & Orchestration.
- Judges may deduct up to 5 points for presentations that significantly exceed the allotted time.